

vLLM

Efficient Memory Management for Large Language Model Serving with PagedAttention

Chenye Wang

Sept 22, 2024

Why do we need vLLM?

The Era of LLMs

- LLM-Powered Services are increasingly important

- However...

- very expensive
- a large number of GPUs

- Solution

- high **throughput** to reduce cost




ChatGPT



GitHub
Copilot

Challenge


- **Batching requests** can improve throughput
 - memory-bound
 - need efficient management for requests' memory -- KV cache
 - huge
 - dynamic
 - unknown lifetime and length
- How do the existing LLM serving systems do?
 - store the KV cache in *contiguous* memory space
 - *pre-allocate* chunk
- What are the problems?



NVIDIA A100 40GB

Memory Fragmentation

- Internal fragmentation
 - actual length can be much shorter
 - inefficiency
 - other requests cannot utilize the part that is currently unused
- External fragmentation
 - different size
 - e.g., buddy allocator



Memory Sharing

- Advanced LLM decoding algorithms generate multiple outputs for one prompt
 - thus sequences can partially share their KV cache
 - e.g., parallel sampling, beam search
- But...
 - KV cache is stored in separate contiguous spaces

How does vLLM do?

PagedAttention

- Inspired by OS
 - *virtual memory and paging techniques*
- Divide KV cache into **blocks**
 - contain a fixed number of tokens
 - not necessarily stored in contiguous space
- Analogy
 - page to block
 - byte to token
 - process to request

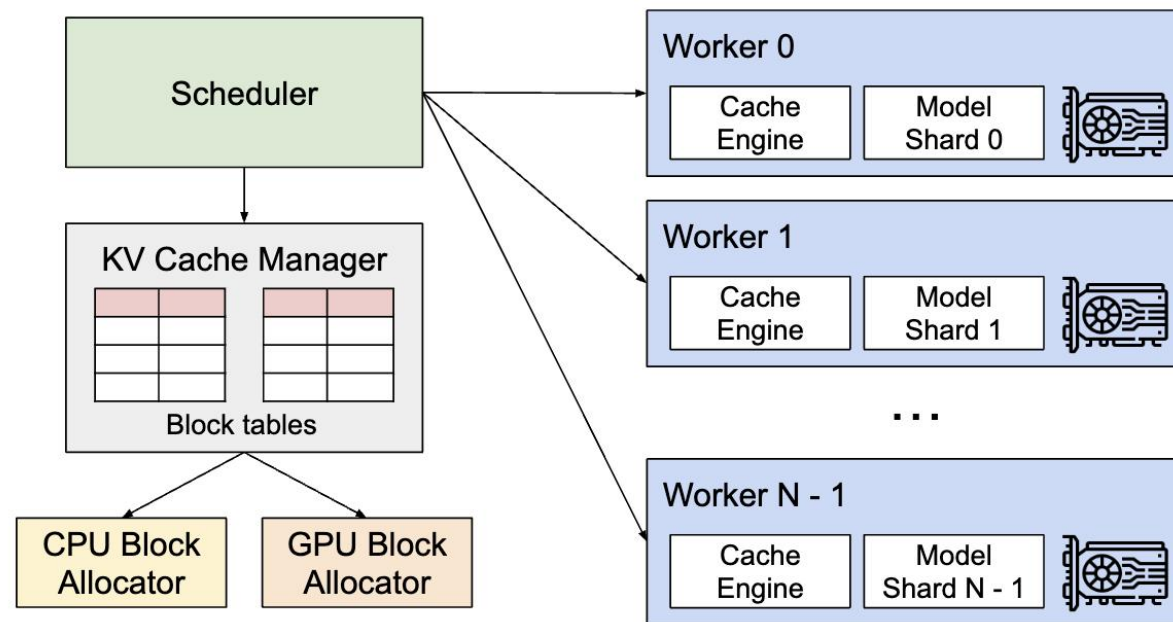



Figure 4. vLLM system overview.

Problems before?

- Internal fragmentation?
 - use small blocks
 - allocate on demand
- External fragmentation?
 - same block size
- memory sharing?
 - at the granularity of a block
 - block mapping



Memory with PagedAttention

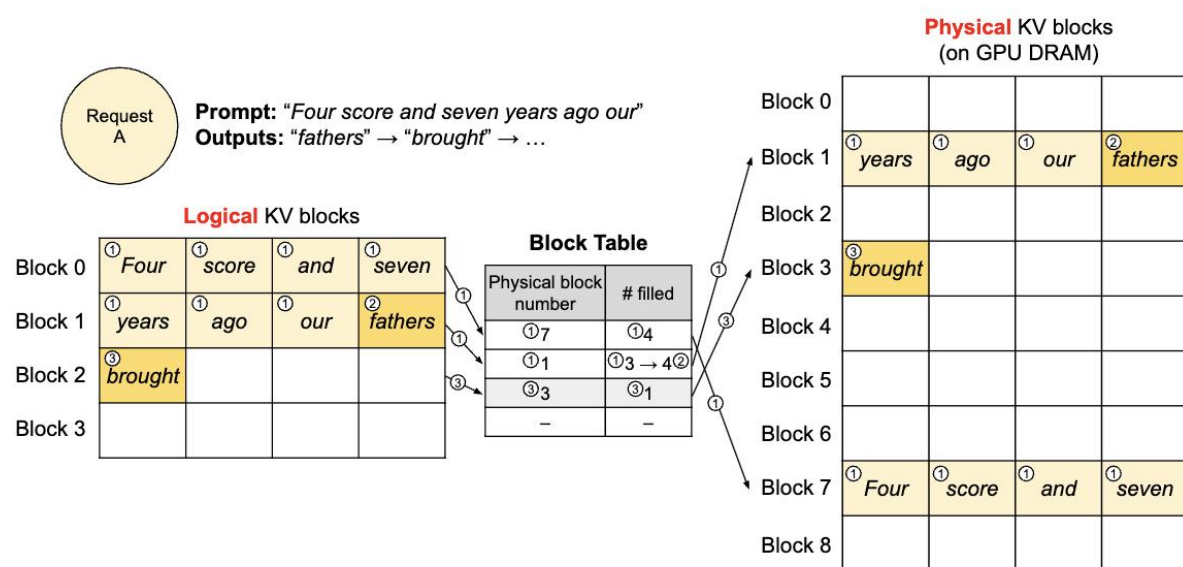


Figure 6. Block table translation in vLLM.

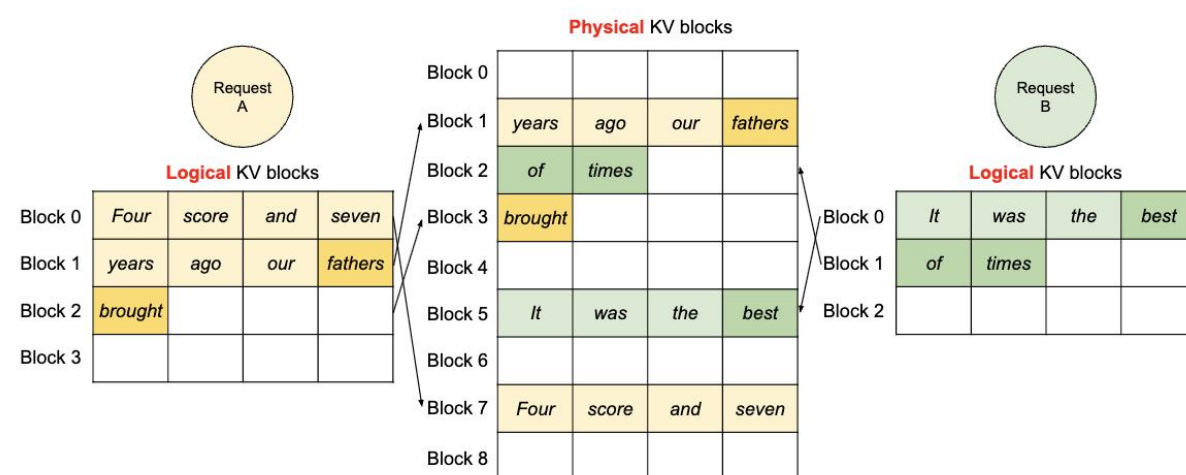


Figure 7. Storing the KV cache of two requests at the same time in vLLM.

Copy-on-Write

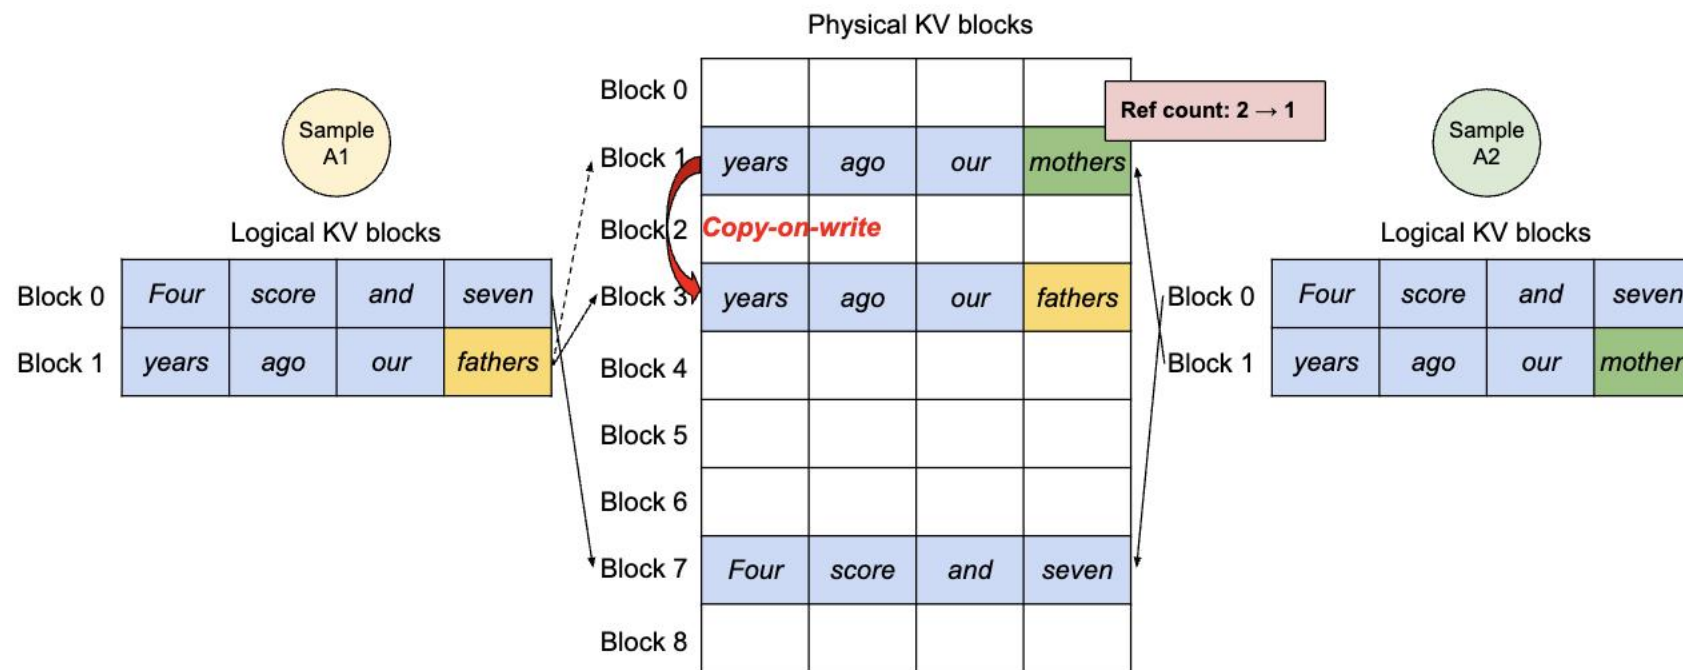


Figure 8. Parallel sampling example.

Other applications

- Beam search
 - similar to *process tree*

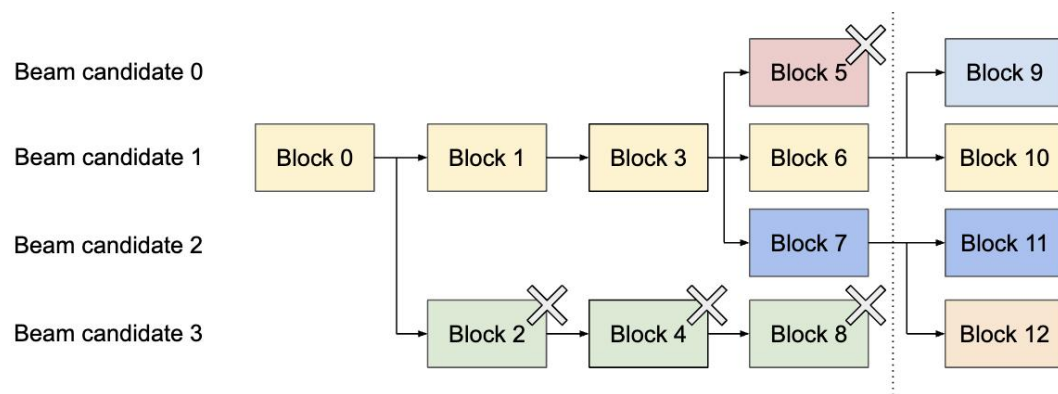


Figure 9. Beam search example.

- Shared prefix
 - similar to *shared library*

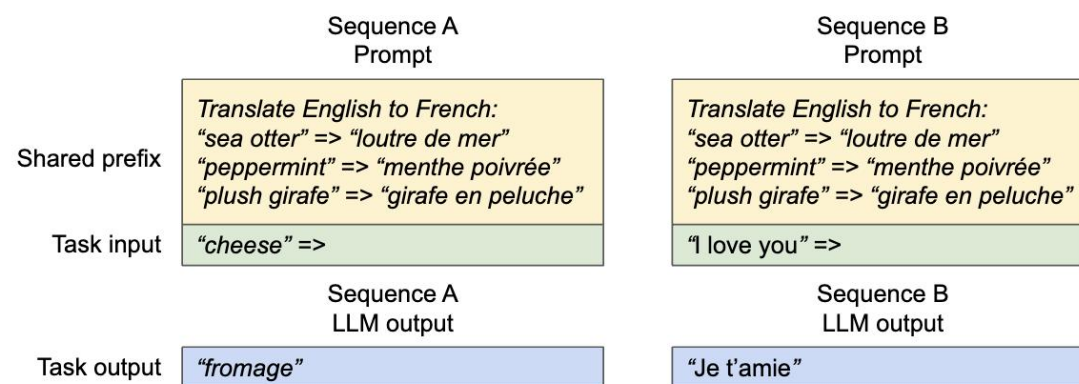


Figure 10. Shared prompt example for machine translation. The examples are adopted from [5].

Kernel-level Optimization

- **Indirect memory access pattern** of PagedAttention is not efficiently supported by existing systems
 - develop custom kernels to optimize it
- fuse operations into kernels
 - KV cache reshape and block write
 - block read and attention operations
 - block copy

Support Various Decoding Algorithms

- vLLM implements various decoding algorithms using three key methods:
 - *fork*
 - create a new sequence from an existing one
 - *append*
 - append a new token to the sequence
 - *free*
 - delete the sequence
- Think of process in OS

What if requests exhaust
GPU memory?

Scheduling and Preemption

- Scheduling
 - FCFS
 - ensure fairness
 - prevent starvation
- Preemption
 - the earliest arrived requests are served first
 - the latest requests are preempted first

Which blocks should vLLM evict?

- the block will be accessed furthest in the future
- **all-or-nothing** eviction policy
 - either evict all or none of the blocks of a sequence
 - multiple sequences within one request are preempted or rescheduled together as a sequence group

How to recover evicted blocks?

- Swapping
 - select a set of sequences to evict, copy their blocks to CPU memory
 - stop accepting new requests until all preempted sequences are completed
 - free blocks of completed requests
 - bring blocks of preempted sequences back
- Recomputation
 - concatenate tokens generated before with original prompt
 - recompute the KV cache

Thinking and Discussion

Tradeoff

- Kernel
 - higher attention kernel latency
 - access the block table
 - execute extra branches
- Block Size
 - too small
 - not fully utilize GPU's parallelism
 - too large
 - internal fragmentation
 - sharing chances decrease
- Recovery mechanisms
 - Swapping
 - depend on the bandwidth between CPU RAM and GPU memory
 - excessive overhead with small block sizes
 - Recomputation
 - depend on computation power of GPU

Discussion for virtual memory and paging

- Applying the same technique to other GPU workloads?
 - stored data shape is static?
 - performance is compute-bound?
- LLM-specific optimizations in applying virtual memory and paging
 - all-or-nothing swap-out policy
 - recomputation for KV cache
 - fusing the GPU kernels for memory access operation

Summary

- Strength
 - flexible memory management
 - near-zero memory waste
 - high hardware utilization and low latency
 - simultaneous processing with different decoding algorithms
- Weakness
 - cannot fundamentally improve inference latency
 - not support sharding the KV cache
 - memory capacity due to large model weights is still the bottleneck

Thanks!

Q & A