

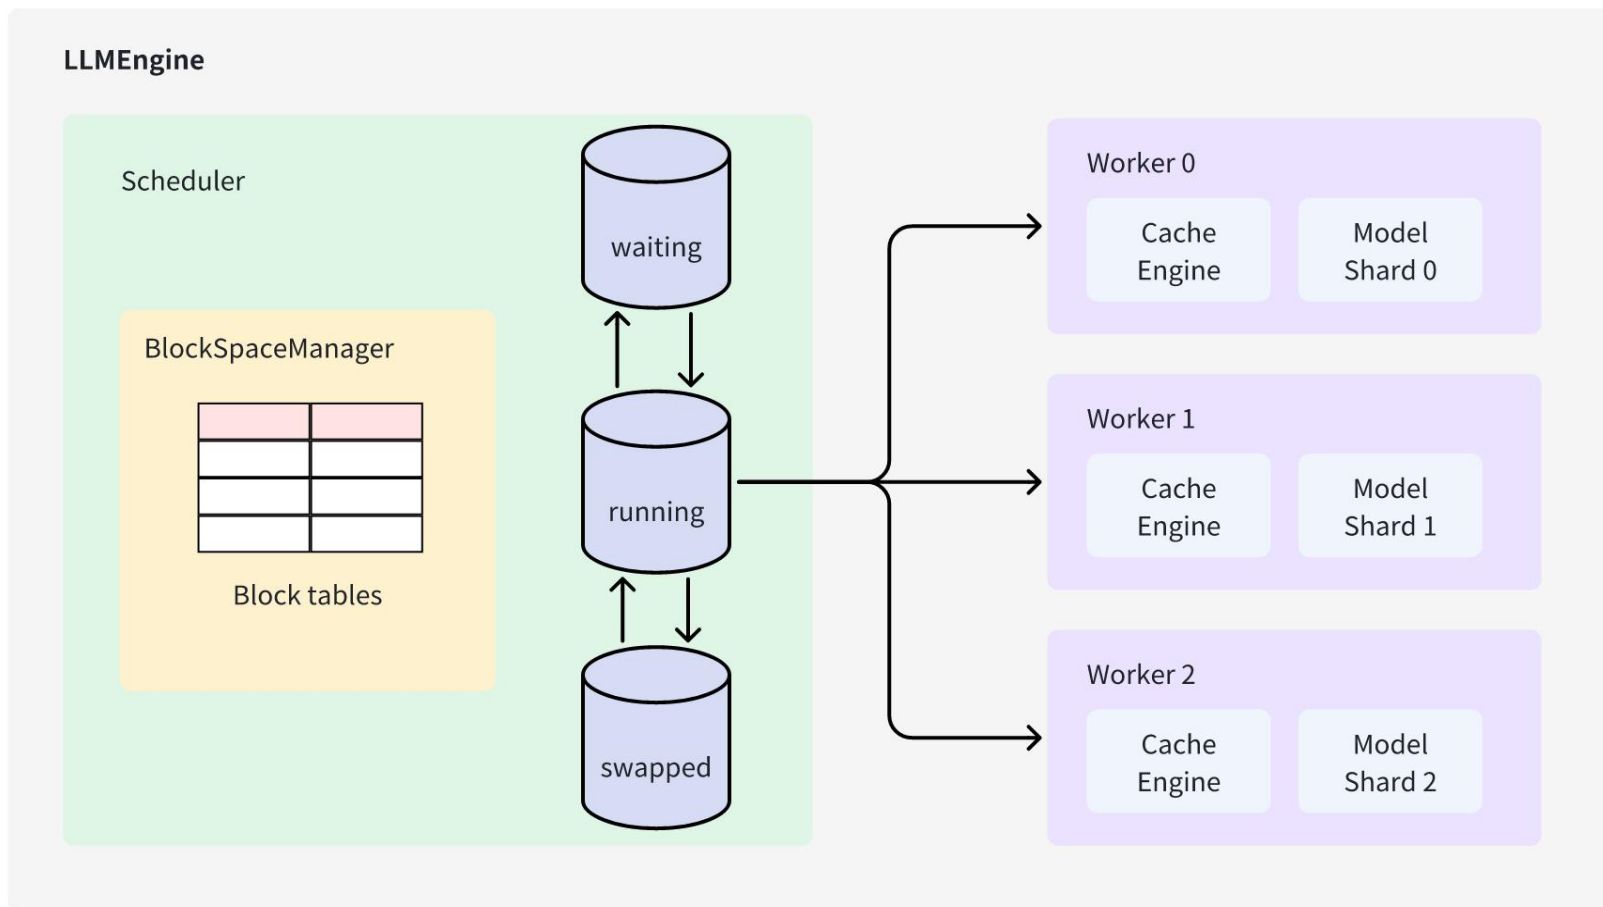
vLLM workflow

Chenye Wang

Oct 29, 2024

Frame

- LLMEngine
 - offline inference
 - online serving
- Scheduler
 - request schedule
- Worker
 - model inference




Scheduler

- Iteration-level strategy
 - reschedule requests after generating a token
 - unfixed batch size ✓
 - different from Orca
 - prefill phase
 - decode phase
- BlockSpaceManager
 - maintain block table
- Queue
 - FCFS


Scheduler

BlockSpaceManager

Block tables

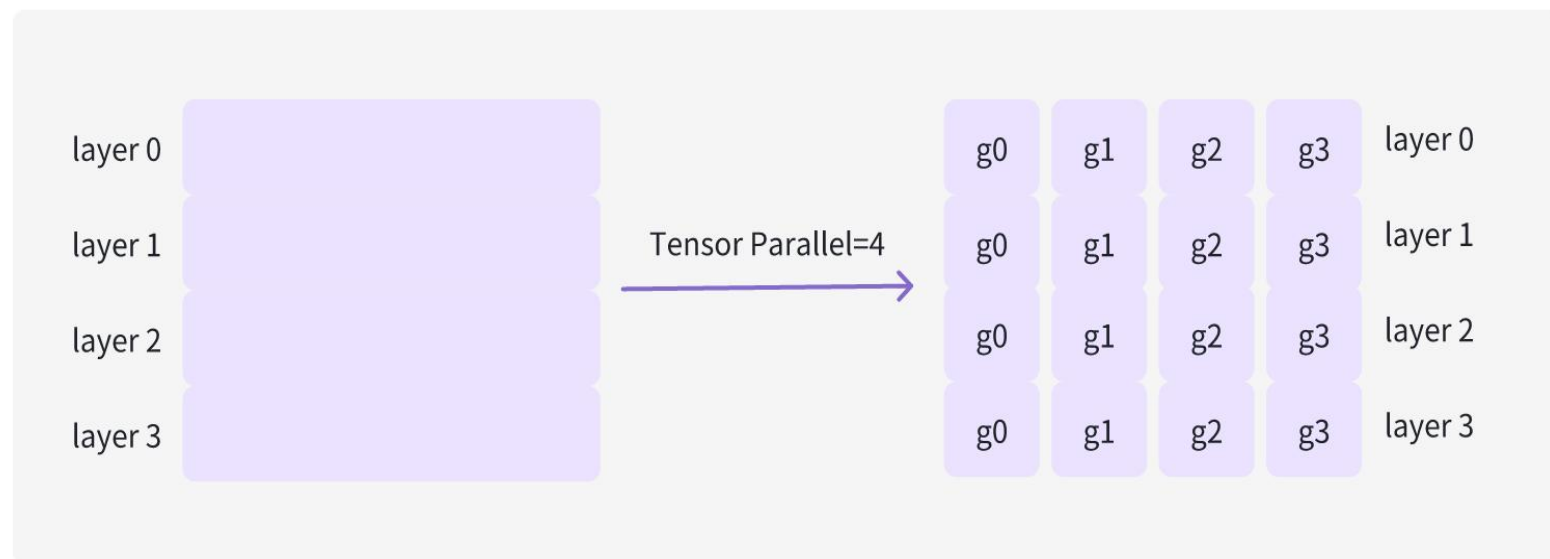
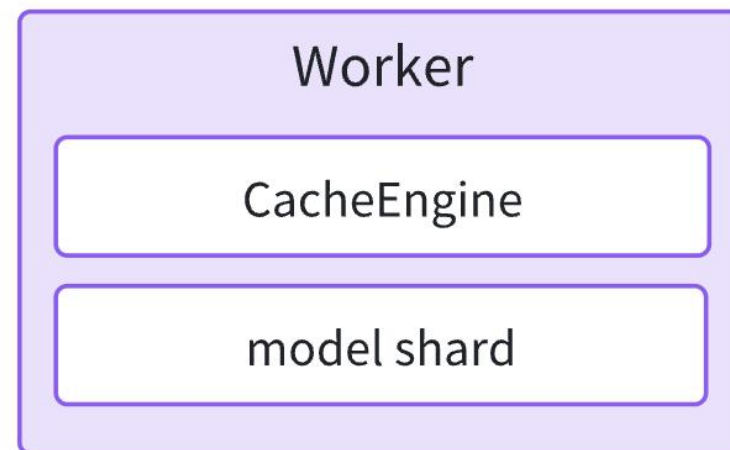


schedule queue



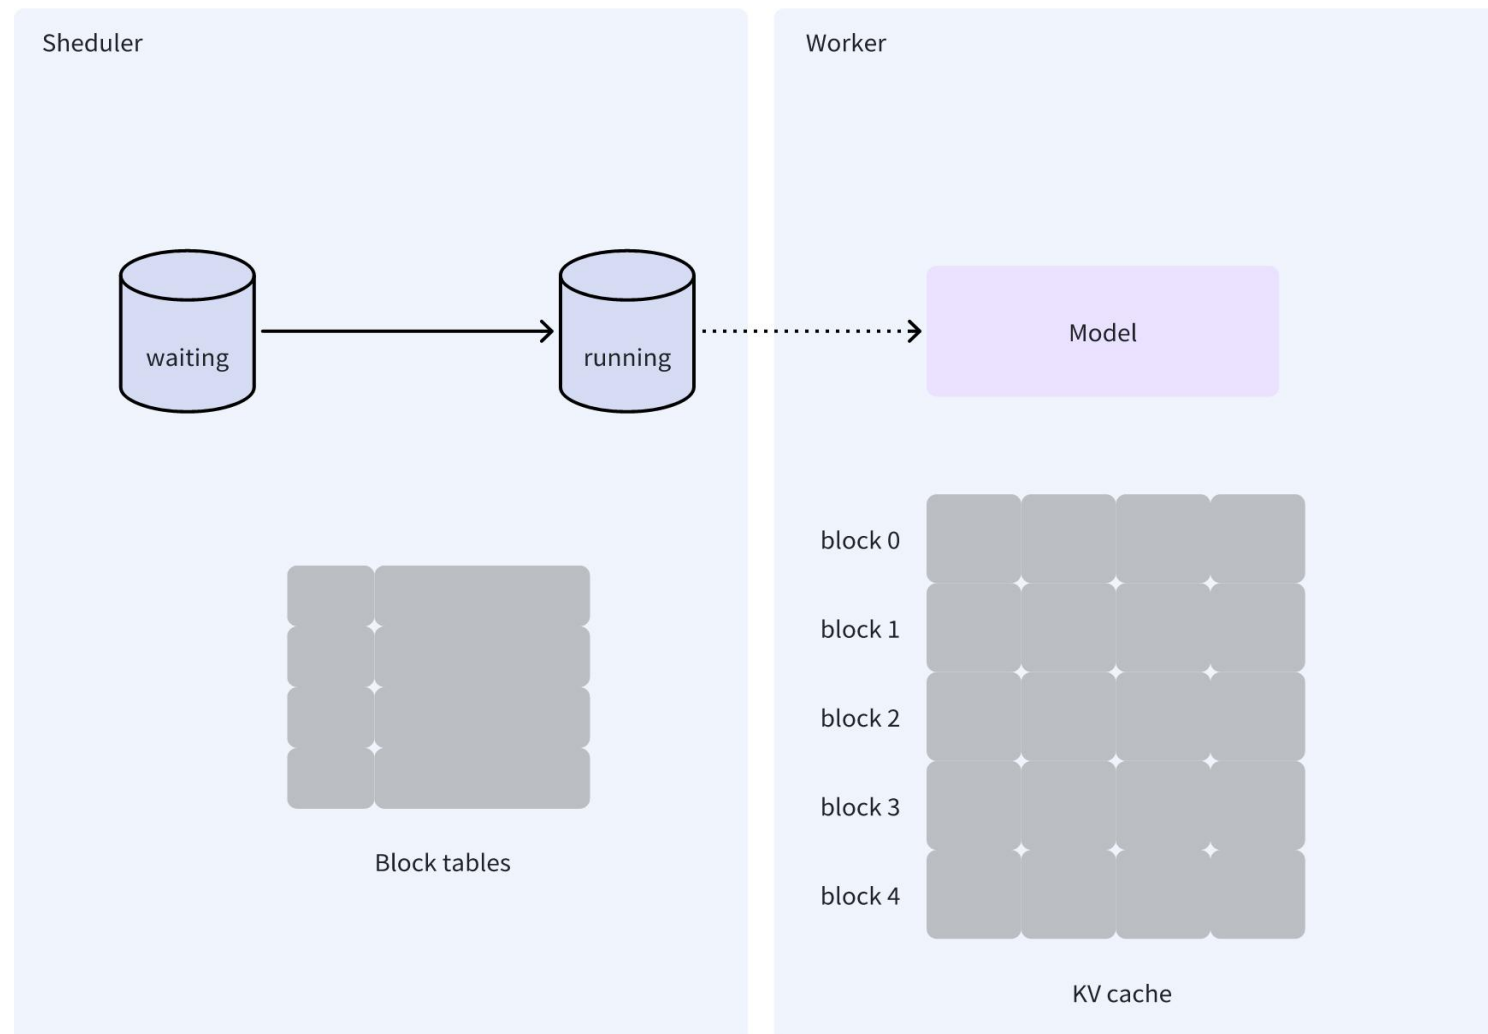
Worker

- Correspond to a GPU
- CacheEngine
 - manage KV cache
- Tensor Parallel
 - distributed



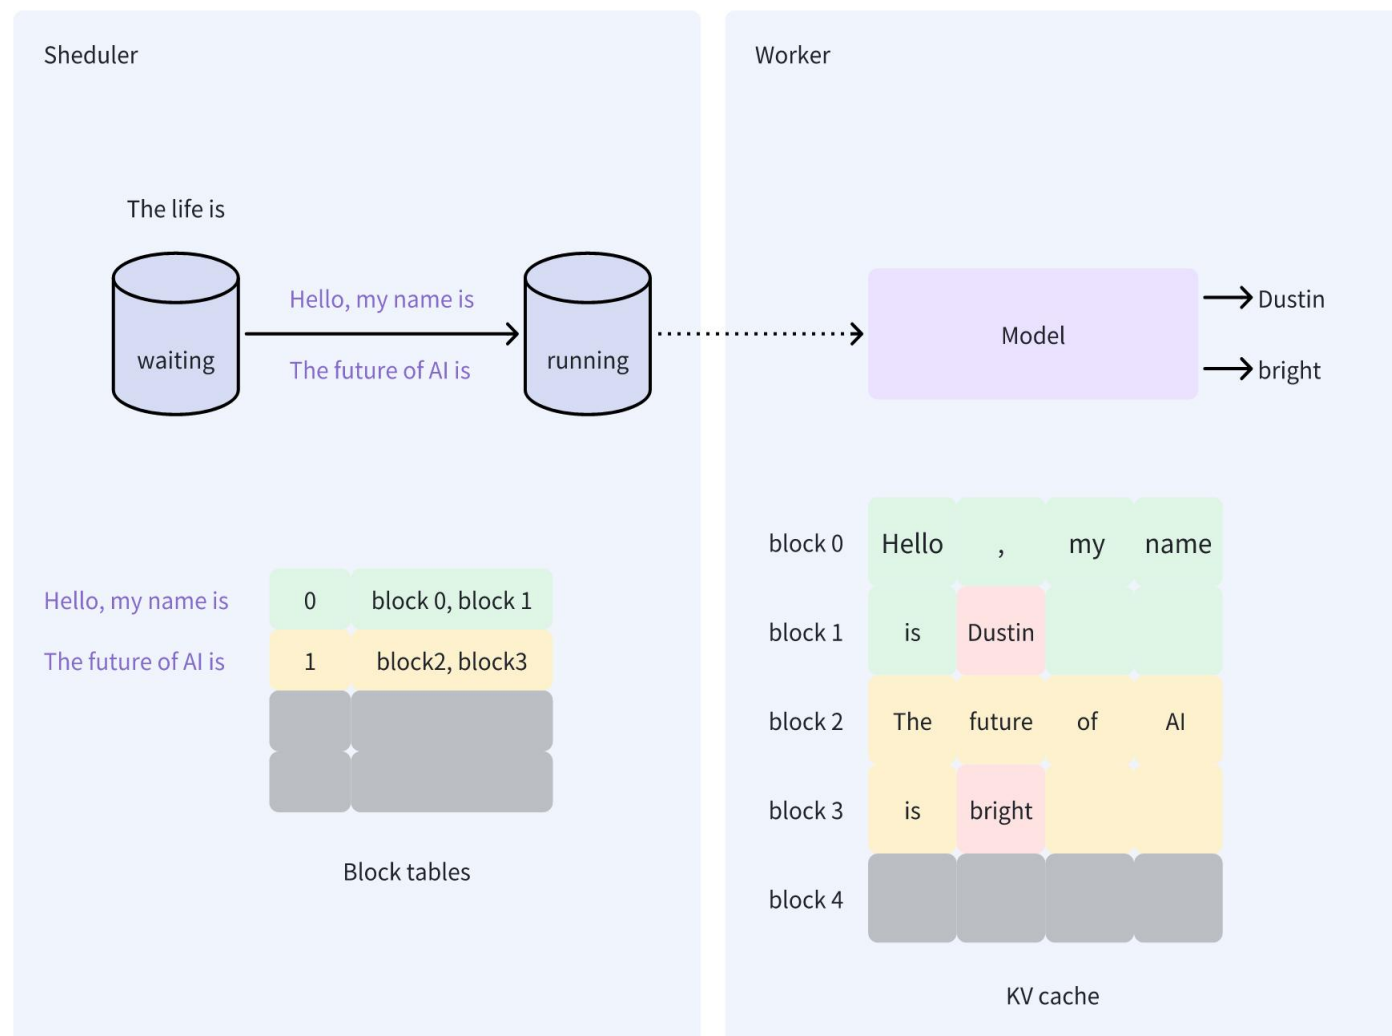
Workflow -- initialization

- Block tables
- Model
- KV cache



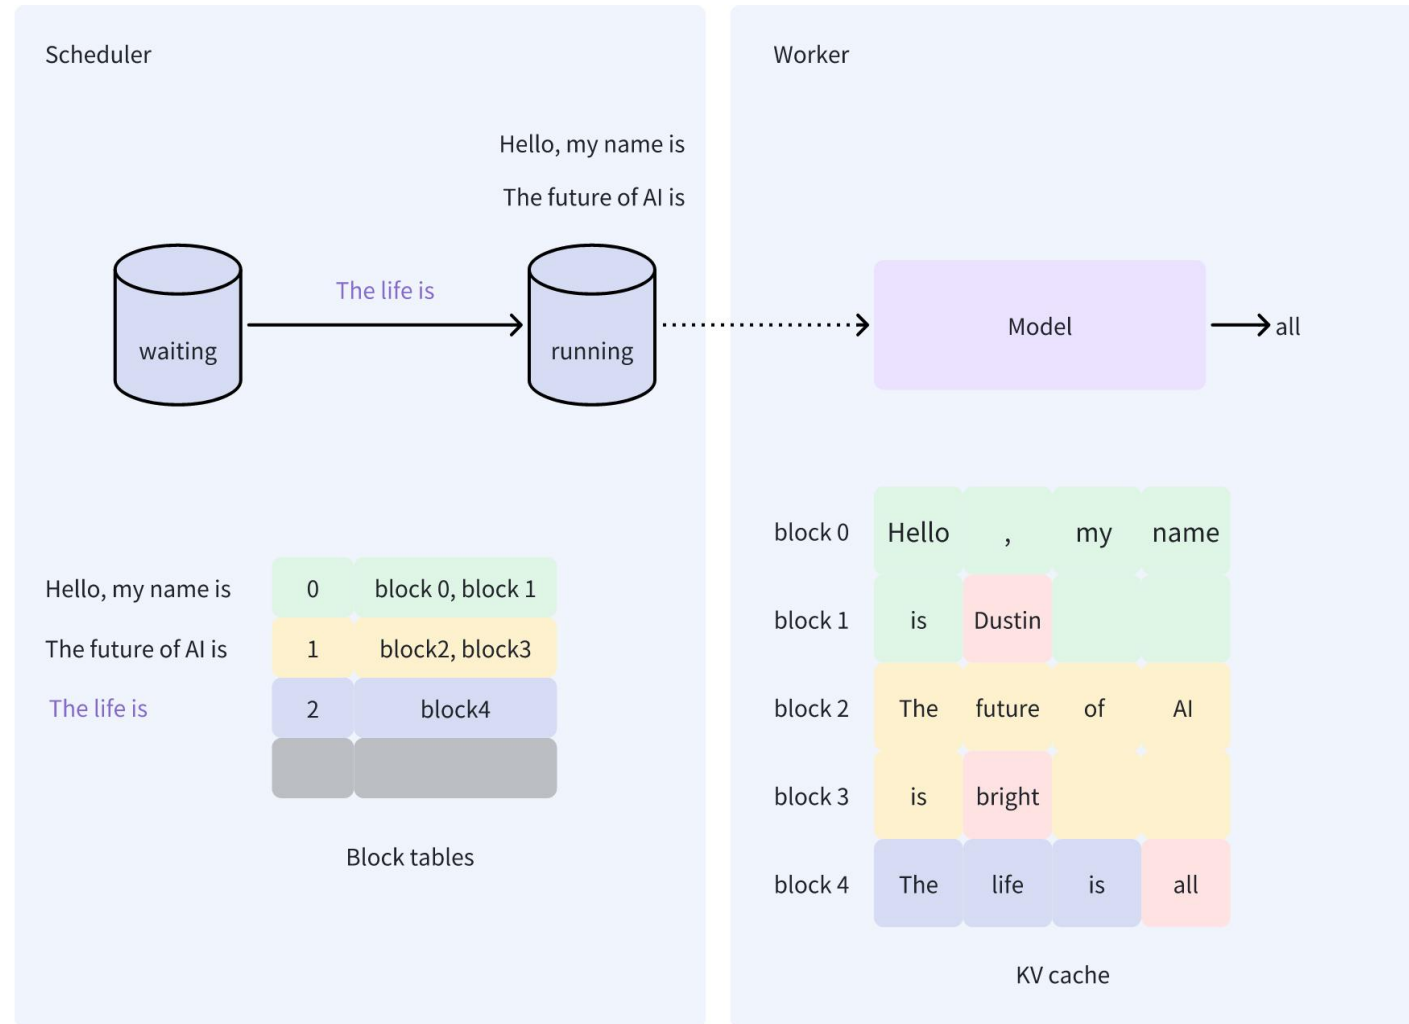
Workflow -- schedule and inference

- FCFS



Workflow -- schedule and inference

- either Prefill
- or Decode



How does vLLM generate output?

- Let's start from a common example, top down

```
prompts = [  
    "Hello, my name is",  
    "The president of the United States is",  
    "The capital of France is",  
    "The future of AI is",  
]  
  
sampling_params = SamplingParams(temperature=0.8, top_p=0.95)  
  
llm = LLM(model="facebook/opt-125m")  
  
outputs = llm.generate(prompts, sampling_params)  
  
for output in outputs:  
    prompt = output.prompt  
    generated_text = output.outputs[0].text  
    print(f"Prompt: {prompt!r}, Generated text: {generated_text!r}")
```


Enter LLM

- Initialize LLMEngine

```
# vllm/entrypoints/llm.py
class LLM:
    def __init__(
        self,
        model: str,
        ...
    ) -> None:
        engine_args = EngineArgs(
            model=model,
            ...
        )
        self.llm_engine = LLMEngine.from_engine_args(
            engine_args, ...)
        ...
```

Generate

- Add all requests, then run engine

```
# vllm/entrypoints/llm.py
class LLM:
    def generate(
        self,
        prompts: ...,
        sampling_params: ...,
        prompt_token_ids: ...,
        use_tqdm: bool = True,
        ...
    ) -> List[RequestOutput]:
        ...
        self._validate_and_add_requests(
            prompts=...,
            params=sampling_params,
            ...)

        outputs = self._run_engine(use_tqdm=use_tqdm)
        return LLMEngine.validate_outputs(outputs, RequestOutput)
```

Run engine

- If has unfinished request, step!

```
# vllm/entrypoints/llm.py
class LLM:
    def _run_engine(
        self, *, use_tqdm: bool
    ) -> List[Union[RequestOutput, EmbeddingRequestOutput]]:
        ...
        while self.llm_engine.has_unfinished_requests():
            step_outputs = self.llm_engine.step()
            for output in step_outputs:
                if output.finished:
                    outputs.append(output)
                ...
            ...
        return sorted(outputs, key=lambda x: int(x.request_id))
```

Enter LLMEngine and ...

